

Roadmap s3:

Cycle 1:

1. Data cleaning & EDA

Missing values, outlier handling

Pandas for eda & plots

Feature engineering-1

Categorical encoding(one-hot, ordinal), scaling and normalisation

2.

Feature engineering-2

Feature selection(correlation, mutual information)

Interaction features

Linear regression

Ordinary least squares

Regularisation(ridge&lasso)

3. Evaluation metrics(mse, r2)

Classification Basics

Logistic Regression theory & sklearn

Metrics: accuracy, ROC-AUC

linear discriminant analysis

4. Decision trees

Tree structure & criteria

5. Random forests

RF hyperparameters & feature importance

k-Nearest Neighbors

Cycle-2:

6. Bagging and boosting

XGBoost, Adaboost

7. support vectors, SVM

8. SVD and related(on TF-IDF, PCA)

9. Unsupervised Learning

k-Means & hierarchical clustering

Dimensionality reduction: PCA, t-SNE

10. Model Validation & Selection

Cross-validation strategies

Learning curves & bias-variance trade-off

Cycle 3:

11. Hyperparameter Tuning & Pipelines

GridSearchCV, RandomizedSearchCV

Pipeline & ColumnTransformer, model serialization

12. Recommender Systems

Collaborative, content-based filtering

Model Explainability

SHAP/LIME intro & demos

13. Perceptron & Linear Classifier

14. PyTorch Foundations

15. Multi-Layer Perceptron (ANN) Scratch Build

MLP architecture, activations (ReLU, sigmoid, tanh)

Backpropagation, vanishing gradients

Cycle-4:

16. continuation to 15th day

17. Feed-Forward Networks in PyTorch

Loss functions, optimizers

18. CNN Basics

Conv mechanics (kernels, stride, padding), pooling, feature maps

19. BatchNorm, Dropout and other functionalities in cnn

20. Intro to rnn & attention